

# Deep Q-Learning



# Recap MDPs


## MDP:

- A MDP is a Markov reward process with decisions. It is an environment in which all states are Markov.
- A Markov Decision Process is a tuple  $\langle S, A, P, R, \gamma \rangle$
- A policy  $\pi$  is a distribution over actions given states:  $\pi(a|s) = P [A_t = a \mid S_t = s]$

## Action - Value Function( $Q(s,a)$ ):

- The action-value function  $q_\pi(s, a)$  is the expected return starting from state  $s$ , taking action  $a$ , and then following policy  $\pi$
- $q_\pi(s, a) = E_\pi [G_t \mid S_t = s, A_t = a]$
- The optimal action-value function  $q^*(s, a)$  is the maximum action-value function over all policies

# Bellman Optimality Equation




$$q_\pi(s, a) = \mathcal{R}_s^a + \gamma \sum \mathcal{P}_{ss'}^a \sum \pi(a'|s') q_\pi(s', a')$$

An optimal policy can be found by maximising over  $q_*(s, a)$ ,

$$\pi_*(a|s) = \begin{cases} 1 & \text{if } a = \underset{a \in \mathcal{A}}{\operatorname{argmax}} q_*(s, a) \\ 0 & \text{otherwise} \end{cases}$$

# Bellman Optimality Equation




$$q_*(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \max_{a'} q_*(s', a')$$

# Dynamic Programming Approach

- Bellman equation gives recursive decomposition
- Value function stores and reuses solutions

## Policy Evaluation:


- Estimate  $q_{\pi}(s,a)$  given a policy using bellman equation
- Start with any random policy



# Dynamic Programming Approach

## Policy Iteration:

- Generate  $\pi \geq \pi'$  by acting greedily according to the bellman optimality equation.
- Guaranteed to converge to the optimal Policy (Proof by Contraction-Mapping)



# Cons

Cannot work with Continuous variables,since the state space is large.

Difficult to work when MDPs are not specified,which is in most cases

# Model based and Model Free approach

## Model-Based:

- Explore environment & learn model,  $T=P(s'|s,a)$  and  $R(s,a)$  everywhere.
- Use policy-evaluation and policy-iteration on the MDP learnt.
- Not feasible in Large state spaces

## Model-Free:

- Rather than learning a model for the environment learn actual state value or action value functions



# Q Learning

- Utility-Sum of discounted rewards in the future
- For all q-states, s,a Compute  $Q_{i+1}(s,a)$  from  $Q_i$  by Bellman backup at s,a. Until  $\max_{s,a} |Q_{i+1}(s,a) - Q_i(s,a)| < \epsilon$
- Bellman Equation

$$Q_{i+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') \left[ R(s, a, s') + \gamma V_i(s') \right]$$

$$Q_{i+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') \left[ R(s, a, s') + \gamma \max_{a'} Q_i(s', a') \right]$$

# Q Learning with Tables

- Consider a 4x4 grid with a start state and a goal state.
- The grid also consists of frozen blocks and some holes
- Design of the grid remains the same, however there is variable wind in action, which might blow the player away to another block when the player takes an action

<i>SFFF</i>	<i>(S: starting point, safe)</i>
<i>FHFH</i>	<i>(F: frozen surface, safe)</i>
<i>FFFH</i>	<i>(H: hole, fall to your doom)</i>
<i>HFFG</i>	<i>(G: goal, where the frisbee is located)</i>

# Q Learning with Tables

- The reward on entering the goal state is +1, and -1 on entering a hole.
- Episode is terminated on entering either of the two states
- Reward is zero in intermediate states

<i>SFFF</i>	<i>(S: starting point, safe)</i>
<i>FHFH</i>	<i>(F: frozen surface, safe)</i>
<i>FFFH</i>	<i>(H: hole, fall to your doom)</i>
<i>HFFG</i>	<i>(G: goal, where the frisbee is located)</i>

# Exploration vs Exploitation

- To decide upon what action is to be performed, a set probability value( $\epsilon$ ) is used
- A randomly generated number is used to determine whether the agent will take random action(Explore) or take the best greedy action(exploit)
- If the random number is above the probability threshold, the optimal action yielding the highest q-value is selected (exploitation).
- Otherwise, a random action is selected (exploration)

# Exponential Moving Average

- Makes recent samples more important


$$\bar{x}_n = \frac{x_n + (1 - \alpha) \cdot x_{n-1} + (1 - \alpha)^2 \cdot x_{n-2} + \dots}{1 + (1 - \alpha) + (1 - \alpha)^2 + \dots}$$

- Forgets about the past (distant past values were wrong anyway)
- Easy to compute from the running average
- Here  $\alpha$  is the learning rate

$$\bar{x}_n = (1 - \alpha) \cdot \bar{x}_{n-1} + \alpha \cdot x_n$$

# A Model Free Approach : Deep Q Learning

- Approximate Q function using a neural network  $f(x, \Theta)$  where  $\Theta$  are learnable parameters of a neural network,  $x$  is input.
- The input to the neural net is the current state of the environment.
- Produce 4 Q-values for each of the action and take the maximum value among them



# Loss Function and Backpropagation

- Do a feedforward pass for the current state  $s$  to get predicted Q-values for all actions.
- Do a feedforward pass for the next state  $s'$  and calculate maximum overall network outputs  $\max_{a'} Q(s', a')$
- Set Q-value target for action to  $r + \gamma \max_{a'} Q(s', a')$
- Update the weights using backpropagation

$$L = \frac{1}{2} \left[ \underbrace{r + \max_{a'} Q(s', a')}_{\text{target}} - \underbrace{Q(s, a)}_{\text{prediction}} \right]^2$$

# Replay Memory


- Approximation of Q-values using non-linear functions is not very stable
- So, during gameplay all the experiences  $\langle s, a, r, s' \rangle$  are stored into a replay memory
- When training the network, random mini batches from the replay memory are used instead of the most recent transition
- Breaks the similarity of subsequent training samples, which otherwise might drive the network into a local minimum
- Makes training task similar to Supervised Learning.




# Implementation On Taxi-V2 and FrozenLake-V0




# Steps Per Episode on Taxi-V2




# Rewards Per Episode on Taxi-V2



# Reward Per Episode on FrozenLake-V0



# Steps Per Episode on FrozenLake-V0



# Future Plans

Self driving car simulator:

- Planning to train an agent to play in a continuous environment by looking at subsequent image frames and taking actions
- Implement this with the help of CNN's

